

cs101aFinal

May 10, 2022

```
[1]: import numpy as np
import pandas as pd
import torch
import torch.nn.functional as F
import torch.optim as optim
import torchvision
import matplotlib.pyplot as plt
from torchvision import datasets, transforms
from torch import nn, optim
from torch.utils.data import Dataset, DataLoader, random_split
from tqdm import tqdm
from PIL import Image as im
```

0, Define the hyper parameters

```
[2]: n_epochs = 100 #Initially 3.
batch_size_train = 16
batch_size_test = 16
```

1, Loading training and testing data

Organizing images First, we need to create a directory structure to hold our images. The directory training2 will hold 10 directories, corresponding to class labels.

This is so we can load the images using torchvision.datasets.ImageFolder, which expects the format described above.

```
[3]: import os

#Make training2 directory
if not os.path.exists('training2'):
    os.mkdir("training2")
#Create the subdirectories corresponding to class labels
for i in range(10):
    if not os.path.exists("training2/{:02d}".format(i+1)):
        os.mkdir("training2/{:02d}".format(i + 1))
```

```
[4]: import shutil

#Iterate through labels.txt, copying images from training into training2 based
↳ on their class label
file = open("labels.txt","r")
for line in file.readlines():
    text = line.strip().split("\t")
    shutil.copy("training/" + text[0], "training2/{:02d}/".format(int(text[1])))
file.close()
```


```
[5]: transform = transforms.Compose([transforms.Grayscale(),
                                    transforms.Resize((32,32)),
                                    transforms.ToTensor(),
                                    transforms.Normalize((0.5,), (0.5,))
                                    ])

dataset = datasets.ImageFolder('training2',transform=transform)
```

```
[6]: def display_img(img,label):
    print(f"Label : {dataset.classes[label]}")
    plt.imshow(img.permute(1,2,0))

#display the first image in the dataset
display_img(*dataset[0])
```

Label : 01



```
[7]: #Splitting the dataset into a testing and training set
trainset, testset = random_split(dataset, [600, 260])
```

Passing the dataset to a dataloader

```
[8]: trainloader = torch.utils.data.DataLoader(trainset,
    ↪ batch_size=batch_size_train, shuffle=True)
testloader = torch.utils.data.DataLoader(testset, batch_size_test=16,
    ↪ shuffle=True)
```

2 Define the model

```
[10]: class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=1, out_channels=32, kernel_size=3,
    ↪ stride=1, padding=1)
        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv2 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3,
    ↪ stride=1, padding=1)
        self.fc1 = nn.Linear(64*8*8, 10)

    # x represents our data
    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.pool(x)
        x = F.relu(self.conv2(x))
        x = self.pool(x)
        x = x.reshape(x.shape[0], -1)
        x = self.fc1(x)

    return x
```

3 Write the training function and the testing function

```
[11]: def test(model, test_loader, device):
    # evaluation, freeze
    model.eval()
    total_num = 0
    total_correct = 0
    with torch.no_grad():
        for _, (data, target) in enumerate(test_loader):

            data = data.to(device)
            target = target.to(device)

            predict_one_hot = model(data)
```

```

_, predict_label = torch.max(predict_one_hot, 1)

total_correct += (predict_label == target).sum().item()
total_num += target.size(0)

return (total_correct / total_num)

```

```

[12]: def train(model, train_loader, test_loader, num_epoch, learning_rate, momentum, device):
    train_losses = []

    # 1, define optimizer
    optimizer = optim.Adam(network.parameters()) #Using default learning rate of 1e-3

    for epoch in tqdm(range(num_epoch)):
        # train the model
        model.train()

        for i, (data, target) in enumerate(train_loader):

            data = data.to(device)
            target = target.to(device)

            optimizer.zero_grad()

            # 2, forward
            output = network(data)

            # 3, calculate the loss
            loss = F.cross_entropy(output, target)

            # 4, backward
            loss.backward()

            optimizer.step()

            # evaluate the accuracy on test data for each epoch
            accuracy = test(model, test_loader, device)
            print('accuracy', accuracy)

    # 5, save model

```

```
[13]: device0 = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
# use cpu if you do not have gpu installed in your computer
network = Net().to(device0)
train(model=network, train_loader=trainloader, test_loader=testloader,
      ↪num_epoch=n_epochs, learning_rate=learning_rate, momentum=momentum,
      ↪device=device0)
```

```
1%|                                     | 1/100 [00:01<02:49, 1.71s/it]
accuracy 0.12692307692307692
2%|                                     | 2/100 [00:03<02:36, 1.59s/it]
accuracy 0.12307692307692308
3%|                                     | 3/100 [00:04<02:33, 1.58s/it]
accuracy 0.25384615384615383
4%|                                     | 4/100 [00:06<02:35, 1.62s/it]
accuracy 0.36153846153846153
5%|                                     | 5/100 [00:08<02:31, 1.59s/it]
accuracy 0.34615384615384615
6%|                                     | 6/100 [00:09<02:33, 1.64s/it]
accuracy 0.36923076923076925
7%|                                     | 7/100 [00:11<02:34, 1.66s/it]
accuracy 0.43846153846153846
8%|                                     | 8/100 [00:13<02:32, 1.66s/it]
accuracy 0.46923076923076923
9%|                                     | 9/100 [00:14<02:29, 1.64s/it]
accuracy 0.4346153846153846
10%|                                    | 10/100 [00:16<02:27, 1.64s/it]
accuracy 0.4807692307692308
11%|                                    | 11/100 [00:17<02:25, 1.64s/it]
accuracy 0.5192307692307693
12%|                                    | 12/100 [00:19<02:27, 1.67s/it]
accuracy 0.4307692307692308
13%|                                    | 13/100 [00:21<02:27, 1.70s/it]
accuracy 0.5576923076923077
14%|                                    | 14/100 [00:23<02:23, 1.67s/it]
```

accuracy 0.5538461538461539	
15%	15/100 [00:24<02:19, 1.64s/it]
accuracy 0.5384615384615384	
16%	16/100 [00:26<02:18, 1.64s/it]
accuracy 0.5423076923076923	
17%	17/100 [00:27<02:17, 1.65s/it]
accuracy 0.5423076923076923	
18%	18/100 [00:29<02:17, 1.67s/it]
accuracy 0.5923076923076923	
19%	19/100 [00:31<02:17, 1.70s/it]
accuracy 0.5846153846153846	
20%	20/100 [00:33<02:13, 1.67s/it]
accuracy 0.5884615384615385	
21%	21/100 [00:34<02:10, 1.65s/it]
accuracy 0.6192307692307693	
22%	22/100 [00:36<02:09, 1.67s/it]
accuracy 0.5923076923076923	
23%	23/100 [00:38<02:08, 1.67s/it]
accuracy 0.6076923076923076	
24%	24/100 [00:39<02:09, 1.70s/it]
accuracy 0.6423076923076924	
25%	25/100 [00:41<02:07, 1.71s/it]
accuracy 0.6076923076923076	
26%	26/100 [00:43<02:10, 1.76s/it]
accuracy 0.6653846153846154	
27%	27/100 [00:45<02:10, 1.79s/it]
accuracy 0.6423076923076924	
28%	28/100 [00:46<02:02, 1.71s/it]
accuracy 0.6576923076923077	
29%	29/100 [00:48<01:57, 1.65s/it]
accuracy 0.6615384615384615	
30%	30/100 [00:50<01:58, 1.69s/it]

accuracy 0.6615384615384615	
31%	31/100 [00:51<01:58, 1.72s/it]
accuracy 0.6576923076923077	
32%	32/100 [00:53<01:59, 1.75s/it]
accuracy 0.6730769230769231	
33%	33/100 [00:55<01:58, 1.77s/it]
accuracy 0.6576923076923077	
34%	34/100 [00:57<01:57, 1.79s/it]
accuracy 0.65	
35%	35/100 [00:59<01:56, 1.80s/it]
accuracy 0.6692307692307692	
36%	36/100 [01:00<01:54, 1.80s/it]
accuracy 0.6692307692307692	
37%	37/100 [01:02<01:53, 1.80s/it]
accuracy 0.676923076923077	
38%	38/100 [01:04<01:52, 1.81s/it]
accuracy 0.6653846153846154	
39%	39/100 [01:06<01:50, 1.81s/it]
accuracy 0.6653846153846154	
40%	40/100 [01:08<01:48, 1.81s/it]
accuracy 0.6807692307692308	
41%	41/100 [01:10<01:47, 1.82s/it]
accuracy 0.6692307692307692	
42%	42/100 [01:11<01:45, 1.82s/it]
accuracy 0.6653846153846154	
43%	43/100 [01:13<01:43, 1.82s/it]
accuracy 0.6653846153846154	
44%	44/100 [01:15<01:42, 1.83s/it]
accuracy 0.6730769230769231	
45%	45/100 [01:17<01:43, 1.88s/it]
accuracy 0.6730769230769231	
46%	46/100 [01:19<01:40, 1.86s/it]

accuracy 0.6692307692307692	
47%	47/100 [01:21<01:36, 1.82s/it]
accuracy 0.6615384615384615	
48%	48/100 [01:22<01:33, 1.80s/it]
accuracy 0.6807692307692308	
49%	49/100 [01:24<01:30, 1.78s/it]
accuracy 0.6961538461538461	
50%	50/100 [01:26<01:28, 1.77s/it]
accuracy 0.6730769230769231	
51%	51/100 [01:28<01:26, 1.76s/it]
accuracy 0.7	
52%	52/100 [01:29<01:23, 1.75s/it]
accuracy 0.6692307692307692	
53%	53/100 [01:31<01:21, 1.74s/it]
accuracy 0.6846153846153846	
54%	54/100 [01:33<01:19, 1.74s/it]
accuracy 0.6730769230769231	
55%	55/100 [01:34<01:18, 1.74s/it]
accuracy 0.676923076923077	
56%	56/100 [01:36<01:16, 1.74s/it]
accuracy 0.6807692307692308	
57%	57/100 [01:38<01:14, 1.74s/it]
accuracy 0.6730769230769231	
58%	58/100 [01:40<01:12, 1.73s/it]
accuracy 0.676923076923077	
59%	59/100 [01:41<01:11, 1.73s/it]
accuracy 0.6923076923076923	
60%	60/100 [01:43<01:09, 1.74s/it]
accuracy 0.6653846153846154	
61%	61/100 [01:45<01:07, 1.73s/it]
accuracy 0.6807692307692308	
62%	62/100 [01:47<01:05, 1.74s/it]

accuracy 0.6923076923076923	
63%	63/100 [01:48<01:05, 1.77s/it]
accuracy 0.6923076923076923	
64%	64/100 [01:50<01:04, 1.78s/it]
accuracy 0.6884615384615385	
65%	65/100 [01:52<01:02, 1.79s/it]
accuracy 0.6807692307692308	
66%	66/100 [01:54<01:01, 1.80s/it]
accuracy 0.7	
67%	67/100 [01:56<00:59, 1.81s/it]
accuracy 0.6807692307692308	
68%	68/100 [01:58<00:57, 1.81s/it]
accuracy 0.676923076923077	
69%	69/100 [01:59<00:56, 1.81s/it]
accuracy 0.6961538461538461	
70%	70/100 [02:01<00:54, 1.82s/it]
accuracy 0.6884615384615385	
71%	71/100 [02:03<00:52, 1.82s/it]
accuracy 0.6961538461538461	
72%	72/100 [02:05<00:50, 1.82s/it]
accuracy 0.6884615384615385	
73%	73/100 [02:07<00:49, 1.82s/it]
accuracy 0.7076923076923077	
74%	74/100 [02:08<00:47, 1.81s/it]
accuracy 0.6961538461538461	
75%	75/100 [02:10<00:44, 1.78s/it]
accuracy 0.6961538461538461	
76%	76/100 [02:12<00:42, 1.76s/it]
accuracy 0.7038461538461539	
77%	77/100 [02:14<00:40, 1.78s/it]
accuracy 0.7038461538461539	
78%	78/100 [02:16<00:39, 1.80s/it]

accuracy 0.7076923076923077
 79%| | 79/100 [02:17<00:37, 1.81s/it]
 accuracy 0.7115384615384616
 80%| | 80/100 [02:19<00:36, 1.80s/it]
 accuracy 0.6923076923076923
 81%| | 81/100 [02:21<00:33, 1.77s/it]
 accuracy 0.6884615384615385
 82%| | 82/100 [02:23<00:31, 1.76s/it]
 accuracy 0.6884615384615385
 83%| | 83/100 [02:24<00:29, 1.76s/it]
 accuracy 0.6923076923076923
 84%| | 84/100 [02:26<00:28, 1.75s/it]
 accuracy 0.7
 85%| | 85/100 [02:28<00:26, 1.75s/it]
 accuracy 0.7076923076923077
 86%| | 86/100 [02:30<00:24, 1.74s/it]
 accuracy 0.7
 87%| | 87/100 [02:31<00:22, 1.72s/it]
 accuracy 0.7038461538461539
 88%| | 88/100 [02:33<00:19, 1.65s/it]
 accuracy 0.6884615384615385
 89%| | 89/100 [02:34<00:17, 1.60s/it]
 accuracy 0.7192307692307692
 90%| | 90/100 [02:36<00:16, 1.67s/it]
 accuracy 0.7076923076923077
 91%| | 91/100 [02:38<00:15, 1.72s/it]
 accuracy 0.7115384615384616
 92%| | 92/100 [02:40<00:14, 1.75s/it]
 accuracy 0.7038461538461539
 93%| | 93/100 [02:42<00:12, 1.78s/it]
 accuracy 0.7153846153846154
 94%| | 94/100 [02:43<00:10, 1.80s/it]

```
accuracy 0.7076923076923077
 95%|          | 95/100 [02:45<00:09,  1.81s/it]
accuracy 0.7038461538461539
 96%|          | 96/100 [02:47<00:07,  1.82s/it]
accuracy 0.7038461538461539
 97%|          | 97/100 [02:49<00:05,  1.83s/it]
accuracy 0.7038461538461539
 98%|          | 98/100 [02:50<00:03,  1.74s/it]
accuracy 0.7
 99%|          | 99/100 [02:52<00:01,  1.67s/it]
accuracy 0.6923076923076923
100%|          | 100/100 [02:53<00:00,  1.74s/it]
accuracy 0.7115384615384616
```

0.0.1 Saving the model

```
[14]: torch.save(network, './Final_Model.pt')
```